

Guidelines — Commit Messages → Developer Role Classification

Submission:

- Submit via a private GitHub repository (can be made public after the competition).
- Include well-structured Jupyter notebooks (.ipynb) for preprocessing, analysis, and modeling.
- Notebooks must contain Markdown explanations of steps and choices.
- Include a README with a brief explanation of your approach and how to reproduce results.
- No API calls allowed.
- Ensure reproducibility: fix random seeds, specify environment, and keep notebooks runnable end-to-end.

LLM Usage Guidelines:

Direct use of pretrained LLMs is not allowed. You cannot simply feed commit messages to GPT/BERT-style models and copy predictions. The focus is on feature engineering, modeling, and reasoning.

Allowed usage:

- You may use pretrained embeddings (e.g., word2vec, fastText, sentence transformers) as part of feature extraction.
- You may use an LLM if you fine-tune or adapt it on this dataset and can demonstrate improvement compared to your baselines.

Disallowed usage:

- Copy-pasting predictions from a closed API (e.g., GPT-4, Claude, Gemini) without training.
- Using LLMs as a hidden oracle to label or 'guess' roles.

Expected Demonstration:

- Show comparisons against simpler baselines (e.g., logistic regression, random forest, BiLSTM).
- Report how fine-tuning the LLM adds value, not just raw usage.

Metrics:

- Macro F1 score
- Precision and Recall
- Confusion Matrix